

Objective of Week 8 Assignment

- Introduce **data poisoning attacks** into the IRIS dataset.
- Evaluate the effect of poisoning levels: **0%, 5%, 10%, 50%**
- Analyze how poisoning impacts **model performance** (Accuracy, Precision, Recall, F1).
- Log all experiments using **MLflow Tracking Server**.
- Generate and compare:
 - Confusion matrices
 - Debug JSON metadata
 - Model artifacts
- Understand how to **mitigate poisoning attacks** in real ML pipelines.

What is Data Poisoning?

- Intentional corruption of training data.
- Goal: reduce model accuracy or cause targeted misclassification.
- Introduced using:
 - Random noise
 - Label flipping
 - Outlier insertion
- Simulated here using **Gaussian noise injection** into a percentage of training samples.

Experiment Setup

Components Used

- **Dataset:** IRIS (Sklearn)
- **Model:** Random Forest Classifier
- **Poisoning Method:** Gaussian noise injection
- **MLflow:**
 - Tracks parameters
 - Logs metrics
 - Stores artifacts (confusion matrix, debug.json, models)
- **VM Setup:** MLflow server running at <http://127.0.0.1:5000>

Step-by-Step Workflow

Steps Performed

1. Setup & start MLflow server in the VM.

Export MLflow tracking URI:

```
export MLFLOW_TRACKING_URI=http://127.0.0.1:5000
```

- 2.

Run the poisoning experiment with multiple levels:

```
python iris_poisoning_mlflow_debug.py --poison-levels 0.0 0.05 0.10 0.50 --noise-std 1.0
```

- 3.
4. Log all metrics, params, and artifacts to MLflow.
5. Export run metrics to CSV ([week8_metrics.csv](#)).
6. Collect and package artifacts from MLflow.
7. Analyze poisoning impact and prepare observations.

Poisoning Logic in Code

Key Components of `iris_poisoning_mlflow_debug.py`

- `poison_data_debug()`
 - Adds random Gaussian noise to selected samples.
- `plot_and_save_cm()`
 - Generates confusion matrix for each poisoning level.
- `run_experiment()`
 - Runs model training + logging for each poisoning level.
- `main()`
 - Loops through poisoning fractions and logs runs to MLflow.

Metrics Observed

Sample Results (based on our run)

| Poison % | # Samples Poisoned | Accuracy | Precision (Macro) | Recall (Macro) | F1 Score (Macro) | Observation |
|----------|--------------------|----------|-------------------|----------------|------------------|--|
| 0% | 0 | 0.90 | 0.9023 | 0.90 | 0.8997 | Baseline (clean data) |
| 5% | 6 | 0.9333 | 0.9333 | 0.9333 | 0.9333 | Slight improvement – noise acts like regularization |
| 10% | 12 | 0.9667 | 0.9697 | 0.9667 | 0.9665 | Best performance — moderate noise increases model robustness |
| 50% | 60 | 0.9333 | 0.9444 | 0.9333 | 0.9326 | Performance drops — heavy poisoning impacts stability |

Confusion Matrices

Confusion Matrices for Each Poisoning Level

- **0% Poisoning**
 - Clean decision boundaries
 - Minimal misclassification
- **5% Poisoning**
 - Still clean — small noise has negligible impact
- **10% Poisoning**
 - Best performing model, confusion matrix almost perfectly diagonal
- **50% Poisoning**
 - More off-diagonal values
 - Model starts confusing overlapping classes due to heavy noise

Observations & Interpretation

Key Learnings from Results

- The model remains **stable** under mild poisoning (5%).
- Moderate poisoning (10%) improves accuracy — **noise acts like a regularizer**.
- Heavy poisoning (50%) begins to distort feature distributions significantly.
- MLflow artifacts confirm:
 - Mean & Std deviation shift heavily after poisoning
 - Confusion matrices become less diagonal at higher noise
 - Debug JSON files show poisoned sample indices and distribution changes
- Poisoning does not always instantly break a model; ML systems must monitor **gradual degradation**.

Mitigation Strategies

Defense Techniques Against Data Poisoning

- **Data Validation Pipelines**
 - Detect abnormal feature distributions
 - Identify outliers or sudden shifts
- **Statistical Drift Detection**
 - Monitor mean, std, KL divergence
 - Compare clean vs new data distributions
- **Robust Training Methods**
 - Noise-tolerant algorithms
 - Median loss functions
 - RANSAC-based approaches
- **Data Provenance Tools**
 - DVC remote storage
 - Audit logs
- **Human-in-the-loop Review**
 - Manually inspect flagged samples
- **Model Monitoring**
 - Track online accuracy changes using MLflow + custom metrics

Files Included

Files Submitted for Week 8

- **iris_poisoning_mlflow_debug.py**
 - Main experiment script
 - Implements poisoning, training, metric logging, and artifact generation
- **week8_metrics.csv**
 - Aggregated table of all experiment runs
 - Used to compare poisoning impact
- **myscript.txt**
 - Commands executed during the VM experimentation and MLflow setup
- **README-week8.md**
 - Documentation explaining the Week 8 workflow, results, and artifacts

The `artifacts_flat/` Folder — What It Contains

The `artifacts_flat/` folder contains **all artifacts extracted from MLflow**, flattened into one directory for easy submission.

It includes:

- **Confusion matrix PNGs**
 - Visual representation of predictions per poisoning level
 - Helps compare model performance qualitatively
- **Debug JSON files**
 - Contains:
 - poisoned sample indices
 - before/after feature stats
 - metrics per run
 - Useful for auditing poisoning impact
- **Model artifacts (`model.joblib`, `MLmodel metadata`, `conda.yaml`)**
 - Exact model used in each run
 - Enables reproducibility

Goal of the `artifacts_flat/` Folder

- To **submit MLflow artifacts** without requiring the entire `mrruns` directory
- To make evaluation easier for instructors
- To show:
 - you logged artifacts correctly
 - you understand debugging metadata
 - you captured model behavior under poisoning

Conclusion

Final Summary

- Successfully simulated **data poisoning attacks** at 0%, 5%, 10%, and 50%.
- Logged complete experiment pipeline using **MLflow**.
- Observed:
 - Improvement at moderate noise
 - Degradation at heavy poisoning
- Generated confusion matrices, debug metadata, and model artifacts.
- Learned the importance of:
 - Data validation
 - Continuous monitoring
 - Artifact tracking
 - Robust model training
- Delivered a reproducible and fully tracked MLOps experiment.